

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import seaborn as sns
# Standard Scaler
from sklearn.preprocessing import StandardScaler
# Robust Scaler
from sklearn.preprocessing import RobustScaler
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
import os
```

```
os.getcwd()
```

```
    '/content'
```

```
dia=pd.read_csv('/content/health care diabetes.csv')
dia.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
dia.shape
```

```
(768, 9)
```

```
dia.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
dia.isnull().sum()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction 0
Age               0
Outcome           0
dtype: int64
```

## 1. Perform Descriptive Analysis

```
dia.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
# Diabetic patient
Positive=dia[dia['Outcome']==1]
Positive.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>0</b>	6	148	72	35	0	33.6	
<b>2</b>	8	183	64	0	0	23.3	
<b>4</b>	0	137	40	35	168	43.1	
<b>6</b>	3	78	50	32	88	31.0	
<b>8</b>	2	197	70	45	543	30.5	

```
Positive.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	
<b>count</b>	268.000000	268.000000	268.000000	268.000000	268.000000	268.000000	
<b>mean</b>	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	
<b>std</b>	3.741239	31.939622	21.491812	17.679711	138.689125	7.262967	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.750000	119.000000	66.000000	0.000000	0.000000	30.800000	
<b>50%</b>	4.000000	140.000000	74.000000	27.000000	0.000000	34.250000	
<b>75%</b>	8.000000	167.000000	82.000000	36.000000	167.250000	38.775000	
<b>max</b>	17.000000	199.000000	114.000000	99.000000	846.000000	67.100000	

```
# Non-diabetic patient
Negative=dia[dia['Outcome']==0]
Negative.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
<b>1</b>	1	85	66	29	0	26.6	
<b>3</b>	1	89	66	23	94	28.1	
<b>5</b>	5	116	74	0	0	25.6	
<b>7</b>	10	115	0	0	0	35.3	
<b>10</b>	4	110	92	0	0	37.6	

```
Negative.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Di:
<b>count</b>	500.000000	500.0000	500.000000	500.000000	500.000000	500.000000	
<b>mean</b>	3.298000	109.9800	68.184000	19.664000	68.792000	30.304200	
<b>std</b>	3.017185	26.1412	18.063075	14.889947	98.865289	7.689855	
<b>min</b>	0.000000	0.0000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	93.0000	62.000000	0.000000	0.000000	25.400000	
<b>50%</b>	2.000000	107.0000	70.000000	21.000000	39.000000	30.050000	
<b>75%</b>	5.000000	125.0000	78.000000	31.000000	105.000000	35.300000	
<b>max</b>	13.000000	197.0000	122.000000	60.000000	744.000000	57.300000	

Non-diabetic dataset description also shows presence of outliers. Pregnancies has mean 3.298000, 75 percentile is 5 and max is 13. This shows outliers. Similarly, Insulin has mean 68.792000, median 39, 75 percentile is 105 and max is 744. This shows outliers.

## 2. Visually explore these variables using histograms:

Glucose, BloodPressure, Skin Thickness, Insulin, BMI


```
dia['Glucose'].value_counts()
```

```

99      17
100     17
111     14
129     14
125     14
..
191      1
177      1
44       1
62       1
190      1
Name: Glucose, Length: 136, dtype: int64

```

```
plt.hist(dia['Glucose'],edgecolor='black',linewidth=1.0)
plt.show()
```




```
dia['BloodPressure'].value_counts().head()
```

```

70      57
74      52
78      45
68      45
72      44
Name: BloodPressure, dtype: int64

```


```
plt.hist(dia['BloodPressure'],edgecolor='black',linewidth=1.0)
plt.show()
```



```
dia['SkinThickness'].value_counts().head()
```

```
0      227
32     31
30     27
27     23
23     22
Name: SkinThickness, dtype: int64
```


```
plt.hist(dia['SkinThickness'],edgecolor='black',linewidth=1.0)
plt.show()
```



```
dia['Insulin'].value_counts().head()
```

```
0      374
105    11
130     9
140     9
120     8
Name: Insulin, dtype: int64
```


```
plt.hist(dia['Insulin'],edgecolor='black',linewidth=1.0)
plt.show()
```



```
dia['BMI'].value_counts().head()
```

```
32.0    13
31.6    12
31.2    12
0.0     11
32.4    10
Name: BMI, dtype: int64
```

```
plt.hist(dia['BMI'],edgecolor='black',linewidth=1.0)
plt.show()
```



Data Exploration:

(1) Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action:

```
dia['Outcome'].value_counts()
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

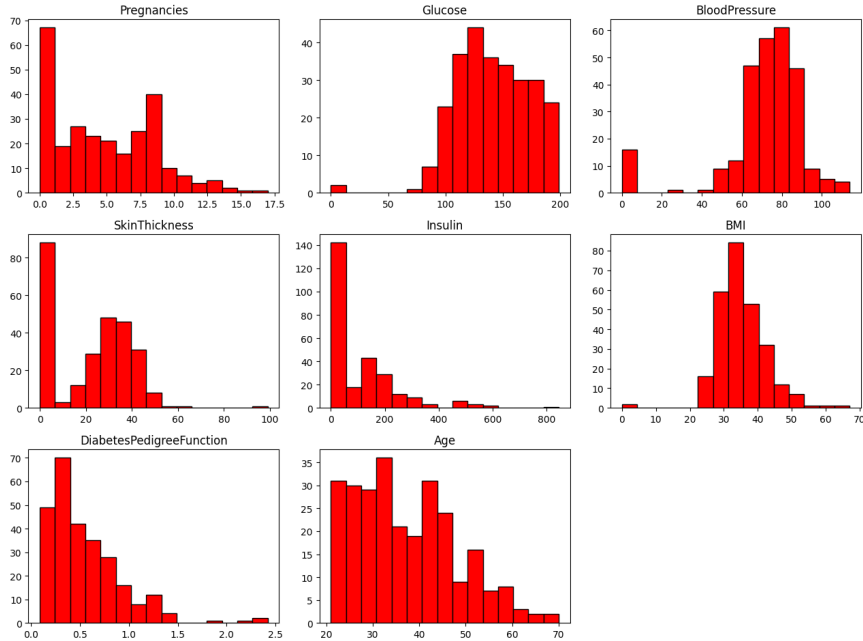
```
Positive=Positive.drop('Outcome',axis=1)
Positive.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
2	8	183	64	0	0	23.3	
4	0	137	40	35	168	43.1	
6	3	78	50	32	88	31.0	
8	2	197	70	45	543	30.5	

```
Positive.hist(bins=15,color='r',edgecolor='black',linewidth=1.0,grid=False)
plt.tight_layout(rect=(0,0,2,2))
plt.suptitle('Outcome 1 (Diabetic Patient)',x=1,y=2.1,fontsize=16)
```

Text(1, 2.1, 'Outcome 1 (Diabetic Patient)')

Outcome 1 (Diabetic Patient)



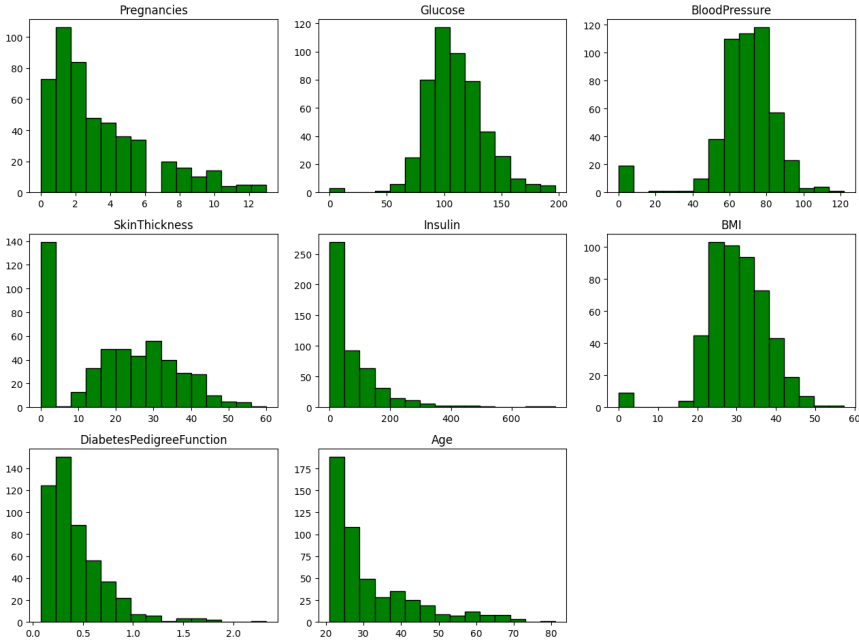
Pregnancies, Insulin, DiabetesPedigreeFunction, Age shows right skew.

```
Negative=Negative.drop('Outcome',axis=1)
Negative.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
1	1	85	66	29	0	26.6	
3	1	89	66	23	94	28.1	
5	5	116	74	0	0	25.6	
7	10	115	0	0	0	35.3	
10	4	110	92	0	0	37.6	

```
Negative.hist(bins=15,color='g',edgecolor='black',linewidth=1.0,grid=False)
plt.tight_layout(rect=(0,0,2,2))
plt.suptitle('Outcome 0 (Non-Diabetic Patient)',x=1,y=2.1,fontsize=16)
```

Text(1, 2.1, 'Outcome 0 (Non-Diabetic Patient)')
Outcome 0 (Non-Diabetic Patient)



Pregnancies, Insulin, DiabetesPedigreeFunction, Age shows right skew.


Scatter Plot

```


BloodPressure=Positive['BloodPressure']
SkinThickness=Positive['SkinThickness']
Glucose=Positive['Glucose']
Insulin=Positive['Insulin']
BMI=Positive['BMI']
Age=Positive['Age']
DiabetesPedigreeFunction=Positive['DiabetesPedigreeFunction']

```

```
g=sns.scatterplot(x= "Glucose" ,y= "BloodPressure",hue='Outcome',data=dia);
```




```
b =sns.scatterplot(x= "BMI" ,y= "Insulin",hue='Outcome',data=dia);
```




```
s=sns.scatterplot(x="SkinThickness",y="Insulin",hue='Outcome',data=dia)
```






```
s=sns.scatterplot(x="Age",y="DiabetesPedigreeFunction",hue='Outcome',data=dia)
```

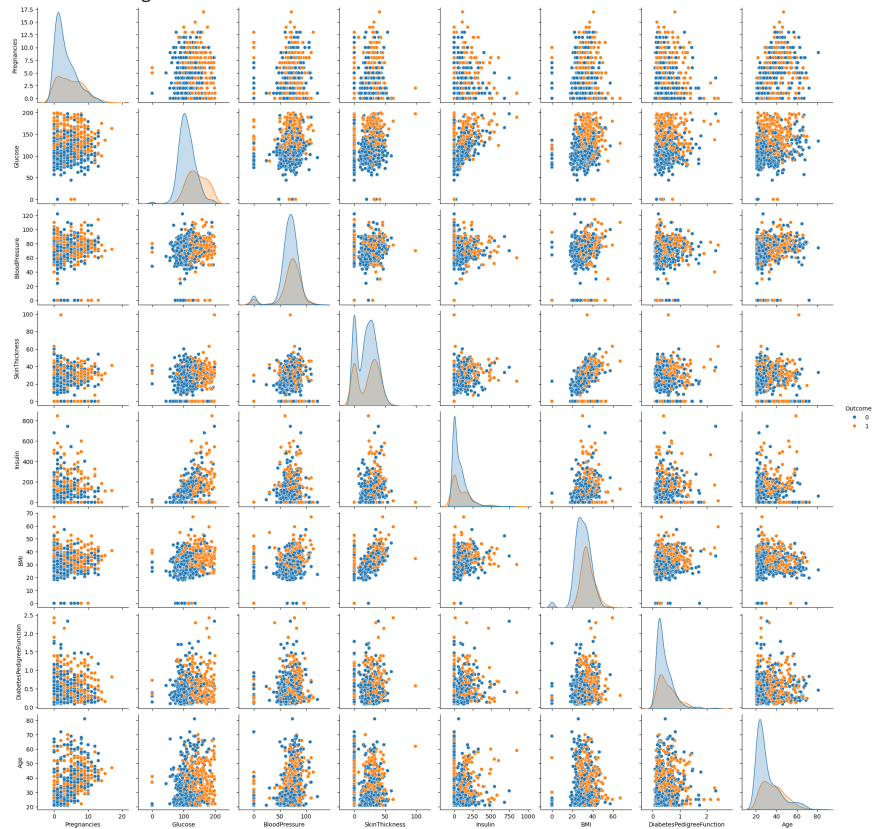


```
s=sns.scatterplot(x="Age",y="BloodPressure",hue='Outcome',data=dia)
```



```
sns.pairplot(dia,hue='Outcome')
```

```
<seaborn.axisgrid.PairGrid at 0x7df607746b00>
```




Pregnancies, Insulin, DiabetesPedigreeFunction and Age are Right skewed. Glucose, BloodPressure, SkinThickness, BMI are almost normally distributed.

```
# Correlation matrix
dia.corr()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
<b>Pregnancies</b>	1.000000	0.129459	0.141282	-0.081672	-0.07353
<b>Glucose</b>	0.129459	1.000000	0.152590	0.057328	0.33135
<b>BloodPressure</b>	0.141282	0.152590	1.000000	0.207371	0.08893
<b>SkinThickness</b>	-0.081672	0.057328	0.207371	1.000000	0.43678
<b>Insulin</b>	-0.073535	0.331357	0.088933	0.436783	1.00000
<b>BMI</b>	0.017683	0.221071	0.281805	0.392573	0.19785
<b>DiabetesPedigreeFunction</b>	-0.033523	0.137337	0.041265	0.183928	0.18507
<b>Age</b>	0.544341	0.263514	0.239528	-0.113970	-0.04216
<b>Outcome</b>	0.221898	0.466581	0.065068	0.074752	0.13054

```
# Heatmap
plt.subplots(figsize=(10,8))
sns.heatmap(dia.corr(),annot=True)
```

<Axes: >




```
dia.skew()
```

```
Pregnancies      0.901674
Glucose           0.173754
BloodPressure     -1.843608
SkinThickness     0.109372
Insulin           2.272251
BMI               -0.428982
DiabetesPedigreeFunction 1.919911
Age               1.129597
Outcome           0.635017
dtype: float64
```

## Normalization

```
dia['Insulin']=np.log(dia['Insulin']+0.1)
dia['Insulin'].hist(bins=15,grid=False)
```

&lt;Axes: &gt;




```
# Apply Robust Scaling to reduce outliers.
rs = RobustScaler(with_centering=True,
with_scaling=True,
quantile_range=(25.0, 75.0),
copy=True)
```

```
dia['Insulin']=rs.fit_transform(dia['Insulin'].values.reshape(-1,1))
```


```
dia.Insulin.hist(bins=15,grid=False)
```

&lt;Axes: &gt;




```
dia['DiabetesPedigreeFunction']=np.log(dia['DiabetesPedigreeFunction'])
dia['DiabetesPedigreeFunction'].hist(bins=15,grid=False)
```

&lt;Axes: &gt;



```
dia['Age']=np.log(dia['Age'])
dia['Age'].hist(bins=15,grid=False)
```

&lt;Axes: &gt;



```
dia.skew()
```

```
Pregnancies      0.901674
Glucose          0.173754
BloodPressure    -1.843608
SkinThickness    0.109372
Insulin          0.005021
BMI              -0.428982
DiabetesPedigreeFunction  0.114178
Age              0.601746
Outcome          0.635017
dtype: float64
```

```
dia['Outcome'].value_counts()
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

Here, the data is imbalanced. So, it is oversampled.

### Oversampling of Data

```
# Synthetic Minority Over-sampling Technique
from imblearn.over_sampling import SMOTE
```

```
smt = SMOTE(sampling_strategy='auto', random_state=9,n_jobs=-1)
X = dia.drop(['Outcome'],axis = 1)
y = dia.Outcome
```

```
X, y = smt.fit_resample(X,y)
```

## Data Modeling

### Logistic Regression

```
# Train-test split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=20)
X_train.shape,X_test.shape,y_train.shape,y_test.shape

((800, 8), (200, 8), (800,), (200,))
```

```
# Create Model
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
# fit the model
model.fit(X_train,y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

```
print(model.score(X_train,y_train))
```

```
0.75375
```

```
print(model.score(X_test,y_test))
```

```
0.74
```

```
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y,model.predict(X))
cm
```

```
array([[377, 123],
       [126, 374]])
```

```
from sklearn.metrics import classification_report
print(classification_report(y,model.predict(X)))
```

```

              precision    recall  f1-score   support

     0       0.75         0.75         0.75         500
     1       0.75         0.75         0.75         500


 accuracy          0.75
 macro avg         0.75
weighted avg         0.75
```

```
#Preparing ROC Curve (Receiver Operating Characteristics Curve)
```

```
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
# predict probabilities
probs = model.predict_proba(X)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(y, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(y, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.833

[&lt;matplotlib.lines.Line2D at 0x7df6006f36d0&gt;]



### Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
model3=DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)
```

```
model3.score(X_train,y_train)
```

```
0.82375
```

```
model3.score(X_test,y_test)
```

```
0.78
```

### Random Forest (Ensemble Technique)

```
from sklearn.ensemble import RandomForestClassifier
model4=RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

```
RandomForestClassifier
RandomForestClassifier(n_estimators=11)
```

```
model4.score(X_train,y_train)
```

```
0.99375
```

```
model4.score(X_test,y_test)
```

```
0.82
```

### KNN

```
from sklearn.neighbors import KNeighborsClassifier
model2=KNeighborsClassifier(n_neighbors=7,metric='minkowski',p=2)
model2.fit(X_train,y_train)
```

```
KNeighborsClassifier
KNeighborsClassifier(n_neighbors=7)
```

```


#Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
# predict probabilities
probs = model2.predict_proba(X)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(y, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(y, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')

```

```

AUC: 0.890
True Positive Rate - [0.    0.262 0.544 0.768 0.886 0.934 0.988 0.998 1.    ], False P
0.28571429 0.14285714 0.    ]
[<matplotlib.lines.Line2D at 0x7df5ff4b97b0>]

```




```

#Precision Recall Curve for Logistic Regression
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(X)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model.predict(X)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y, probs)
# calculate F1 score
f1 = f1_score(y, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```




f1=0.750 auc=0.818 ap=0.819  
 [matplotlib.lines.Line2D at 0x7df5ff522950]



```
#Precision Recall Curve for KNN
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(X)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(X)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(y, probs)
# calculate F1 score
f1 = f1_score(y, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(y, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.820 auc=0.885 ap=0.856  
 [matplotlib.lines.Line2D at 0x7df5feb9be20]



```
#Precision Recall Curve for Decision Tree Classifier
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
```